

StudyOS Thesis-Finder

Better-prepared students, fewer misdirected inquiries

StudyOS Team

July 2026

Department of Computer Science, University of Tübingen

The problem, from both sides

Students

- Spend weeks figuring out which chairs even work on their interests
- See 4–5 listed topics, assume that is the whole lab, and never reach out
- Send generic cold emails to the wrong chairs

Professors

- Receive far more inquiries than they can supervise
(one chair: 50+ inquiries for 20 slots)
- Spend time writing rejections they regret
- Get low-signal first contacts that miss the chair's actual research

Both sides lose to the same failure: low-signal first contact.

What 27 professors in this department told us

We interviewed 27 professors before building anything. The verdict on a classic *topic-listing platform*:

- ~50% do not want a central platform — either already overrun with applicants, or convinced topics must be *co-created in conversation*, not browsed from a list
- ~35% would try a tool — but only if zero-friction, with no reminder emails, and stable for years
- ~15% recruit through other channels entirely

“Profs never update the topics. We’ve had many attempts in the department already.”

A topic platform was tried here around 2022 and failed — for **incentive** reasons, not UI reasons.

Our conclusion: build for students, not for chairs

A tool that requires **zero professor participation** to deliver value.

The agent helps a student to:

1. **Formulate** — turn vague interests, coursework, and skills into a structured competency profile
2. **Discover** — find which chairs and researchers actually work on relevant problems, based on *public* data: publications, team pages, research-area statements
3. **Prepare** — draft a specific, high-signal first contact: who they are, what they can do, why *this* chair

Chairs remain passive data sources by default; their own websites stay the source of truth.


How it works: a pipeline of agent skills

No app, no accounts, no server. The student runs a standard AI coding/chat agent; our package is a set of portable *skills* plus curated data.



- Input can be as thin as *"I like deep learning and robotics"* — the agent interviews the student to fill the gaps
- Every recommendation cites its evidence: papers, people, group pages
- The output coaches students toward a *conversation* with the chair, not topic shopping

The data layer: a curated, public researcher tree



- Built only from **already-public** sources: OpenAlex, publication records, official team pages
- Stored as plain, reviewable Markdown — every entry carries provenance, dates, and uncertainty flags
- Refreshed by an automated monthly job as a reviewable change request; manual corrections are never silently overwritten
- Validated against hand-verified ground truth for pilot chairs, then scaled to all 47 professors

How this answers the concerns you raised

You told us...	So we...
"Don't build a system we have to use"	built a student tool; chairs do nothing by default
"Read our existing websites instead"	ingest public pages and publications — no data entry
"Topics are co-created, not browsed"	surface <i>research areas and people</i> , and coach students toward conversation
"Too many misdirected inquiries"	better-matched, better-prepared students $\Rightarrow$ fewer cold, low-signal emails
"Will it still exist in 3 years?"	no chair investment to lose; your websites remain the source of truth

- **Student data never leaves the session.** The profile (interests, transcripts, constraints) exists only in the student's own agent conversation — nothing is stored or transmitted by us.
- **Only public academic data is bundled:** names, roles, publications, and research areas as published by chairs and OpenAlex. We document this policy explicitly with GDPR in mind.
- **Everything is inspectable.** Data and skills are plain text in a public repository; any chair can review — or ask to correct — what is said about their group.
- **No fabricated numbers.** A hard project rule: no placeholder metrics, no invented team sizes or citation counts.

Where we are, and what happens next

Now — Phase 1: get the data right first

- Pipeline piloted on 3 chairs with hand-verified rosters as ground truth
- Gate before anything else: $\geq 90\%$ recall on pilot rosters, full referential integrity, reproducible golden record

Then — Phase 2: optimize the advising behavior

- Scale the researcher tree to all 47 professors
- Systematic evaluations: shallow profiles, missing information, degree-program and thesis-scope constraints (B.Sc. vs. M.Sc.)

Ordering principle: matching cannot be trusted while the underlying data are incomplete or stale — so data quality is gated before behavior tuning.

What we ask of you (very little)

1. **Nothing, by default.** The tool works from your public pages and publications as they are.
2. **Optionally:** a handful of research-area keywords — data that rarely needs revision. No logins, no dashboards, and a hard rule of **no reminder emails**.
3. **Feedback:** if inquiries you receive start arriving better-prepared, we would like to hear it — that is our success metric, not platform adoption.

Students get better discovery. You get fewer but better inquiries.
Nobody is forced to maintain anything.

Thank you — questions?